

Occupancy Analyzer

Global Software Development Project

Christoffer Wilken Pagaard
Jacob Romme Rasmussen
Theresa Brandt von Fackh
Tomas Miseikis



IT University
of Copenhagen

15. December 2013

Abstract

Write abstract text...

Contents

1	Introduction	5
1.1	Context	5
1.2	Problem	6
1.3	Related Work	6
1.4	Approach	8
1.5	Report Structure	8
2	Analysis	9
2.1	Image Analysis	9
2.2	Prediction	10
2.3	Communication Model	12
2.3.1	Central Server	12
3	Design	13
3.1	System Overview	13
3.2	Web Cameras	14
3.3	Raspberry Pi Computers	14
3.3.1	Object Extraction	15
3.3.1.1	Simple Background Subtraction Approach	15
3.3.1.2	Running Average Approach	17
3.3.2	Object Detection	18

3.3.3	Object Differentiation	20
3.4	The Server	21
3.4.1	Data Input API	21
3.4.2	Data Output API	21
3.4.3	Storage	21
3.5	Prediction	21
3.5.1	Gathering statistical data	21
3.5.2	Custom prediction model based on HMM	22
3.6	Android Application	23
4	Implementation	27
4.1	Image Processing	27
4.1.1	Library	27
4.1.2	Programming Language	27
4.1.2.1	Performance Comparison Between Java and Python	28
4.2	The Server	28
5	Evaluation	31
5.1	Verification	31
5.2	Benchmarks	31
6	Collaboration	32
6.1	Theory	32
6.1.1	Project management	32
6.1.2	Global collaboration	32
6.2	Background information of the Project Team	33
6.2.1	Team ITU	33
6.2.2	Team Kenya	34

6.3	Project Management	35
6.3.1	Approach	35
6.3.2	Project Organization	35
6.3.2.1	Timeline	35
6.3.2.2	Project Team and Roles	35
6.3.3	Project Management Tools and Methods	36
6.3.3.1	Member profiles	36
6.3.3.2	Skill-/Preference- and Motivation-Sheet	36
6.3.3.3	"Out-of-Office"-Calendar	36
6.3.3.4	Milestone plan	37
6.3.3.5	Internal weekly meeting (Jour-Fixe)	38
6.3.3.6	Global weekly meeting (Jour-Fixe)	39
6.3.3.7	Meeting reports	39
6.3.3.8	Time recording	40
6.3.3.9	Introduction and Kick-Off meeting	40
6.3.3.10	Surveys	40
6.3.4	Collaboration Tools	40
6.3.4.1	Communication	40
6.3.4.2	Sharing	41
6.3.4.3	Time recording	41
6.4	Project Issues	41
6.4.1	Collaboration Issues	41
6.5	Hypothetical Scenarios	41
7	Discussion	42
8	Conclusion	43

1 | Introduction

1.1 Context

Smart use of energy resources is an ongoing topic these days. The reduction of expenses is mostly the biggest driving factor for companies. But also the debate around climate change brings new legislation to reduce the waste of energy resources, whose production is damaging to the environment and future generations. The IT University of Copenhagen (ITU) has an interest in producing an occupancy model for commercial buildings, like the ITU building, to detect where energy resources are needed and where they can be saved. Energy resources are needed for e.g. lighting and heat-regulating systems, which are relevant for occupants in a commercial building. With the detected occupancy data the ITU can predict occupancy and develop concepts for a smart use of energy resources in commercial buildings.

The Strathmore University in Kenya has also an interest in building up an occupancy model, but mainly for surveillance reasons. Surveillance can be used for several purposes like traffic monitoring, public safety and facilities surveillance. An IT-based surveillance system can automatically analyze the scene without the use of human resources. By analyzing the scene the detection of occupancy is a major part. Moreover a real-time prediction model on top of the occupancy data can be used to prevent criminal activity by triggering alarms or other surveillance systems.

Currently there is no existing infrastructure to build up an occupancy model in the Strathmore University or the ITU building. Both universities want a solution for an occupancy analyzer based on Raspberry Pi computers due to the minimal consumption of computational and monetary resources. Furthermore, Starthmore University requests for an Android application, which represents a live-feed of the occupants in a monitored room.

A group of students from both universities have to collaborate to come up with a solution for an occupancy analyzer, which can satisfy the needs of both university interests. Ideally a product should be developed, which can be adapted to fit the needs of one or the other university. Furthermore, a collaboration project is mandatory for the student group from ITU, in which they have to face the challenges of global collaboration, navigate compromises and come up with a solution.

This report contains the product result, design of the product, details of the project work and the learning outcomes, which were achieved in the project with the globally distributed team from the perspective of the ITU students. The project team consists of international students located in Nairobi, Kenya (East African Time) and Copenhagen, Denmark (Central European Time).

1.2 Problem

The content of this project is to build an occupancy analyzer, which detects people in a room or corridor and predicts their movement. The occupancy analyzer has to be based on Raspberry Pi computers, which come with computational restrictions, and web cameras. A solution for the right architecture and programming languages has to be found, which can deal with those limited resources.

Due to the usage of web cameras, a visual detection of people has to be made. Analyzing images by detecting people - which are moving objects and not part of the room - is one major challenge to face. Visual conditions of the image can change as a result of daylight. For example, dynamic lighting and moving shadows should not influence the detection of people. To best capture occupants and the requirement to represent the occupants on an Android application, the position of web cameras becomes an important factor. The differentiation between multiple people, as well as between people and the setting of the room, is important for the quality of the detection. Only if reliable data about occupancy exists, can the data be used to construct a reliable prediction model usable for future concepts and projects.

The question of how the collected data can be used, has to be considered. Building prediction models, which relies on historical data and real-time data, is another requirement, which the project dealt with. Decisions on what kind of prediction for a meaningful application - like the one mentioned in paragraph 2 in Section 1.1 - and how the data will be stored and processed have to be made.

Besides the design and implementation of an occupancy analyzer, another task is the collaboration of students from two different universities. Cultural differences, difference in time, spatial distance and locally related influences have to be overcome. Different perspectives have to be combined and compromises have to be made.

1.3 Related Work

Background subtraction is an extremely important concept in surveillance applications, since this is the first step in detecting people and tracking their movement in monitored area. Y. Benetzeth, P.-M. Jodoin, B. Emile, H. Laurent & C. Rosenberger [1], in their work *Compar-*

tive Study of Background Subtraction Algorithms, mention a basic motion detection approach by using a simple background subtraction technique. In this technique, the background image is taken when there is a total absence of motion in the monitored area. This background image is then subtracted from every newly taken image of the area in order to detect motion. This is possible since after subtraction the resulting image will either be almost totally black - meaning no motion occurred - or the image will have some resulting bright contours of detected differences.

We applied this simple technique in the beginning of the project for the initial experiments and studies of object detection and tracking. These experiments and studies provided us with some knowledge on how basic object detection works and how we can apply it for our specific problem. Moreover, we have found that this technique can work well in static environments, but falls short in dynamic ones due to illumination changes and background modifications in the monitored area.

Aditi Jog & Shirish Halbe [2], in their paper *Multiple Object Tracking using CAMshift Algorithm in Open CV*, present a better way of building object detection and tracking application than simply using basic background subtraction approach. For their implementation they used CAMshift¹ algorithm, which is capable of tracking multiple moving objects in video streams. In this algorithm, an approximate background image of the monitored area is gradually built up by smoothing or blurring out the motion areas. This is done by using several images of monitored area taken prior, and performing arithmetic averaging on them. One can then simply use this built up background image for detecting changes in the monitored area.

This approach applies to the first part of our problem, namely detecting and tracking people in the monitored area. We can build upon this approach as the implementation of it is quite simple yet flexible and adaptive. Moreover, it eliminates the need to worry about lighting fluctuations or other unpredictable background changes that may corrupt surveillance of the monitored area.

Gellert & Vintan [3], in their work *Person Movement Prediction Using Hidden Markov Models*, presents the use of hidden Markov models to anticipate the next movement of a given person. It is used inside an office building where it is possible to predict the next room given any room based on the history of rooms visited by a person moving within the building. The configuration of the model is optimised by evaluating movement sequences of real persons within an office building.

The concept of calculating the probability, and thereby making a prediction, of the next room given any room can be related to our scenario where we predict the next occupied section of the image given an occupant in any section. This information is used in the process of predicting where an occupant will exit the room.

¹Continuously Adaptive Mean Shift.

Ashbrook & Starner [4] demonstrated how locations of significance can be incorporated into a predictive model in the paper *Using GPS to Learn Significant Locations and Predict Movement Across Multiple Users*. They cluster location data from a GPS over an extended period of time and infer meaningful locations. These locations are then incorporated into Markov models where it is used to predict a person's future location given any current location. The idea of using heavily visited areas, that is *locations of significance*, to perform predictions is transferred to our scenario, where each area of the image holds a value reflecting the amount of activity within the area. This value plays an important role in the calculations of probabilities and thereby the final predictions.

1.4 Approach

1.5 Report Structure

2 | Analysis

2.1 Image Analysis

There are numerous factors that must be considered when designing a surveillance system.

Position of the camera. Camera's type, position and angle may differ a lot, depending on the requirements of the surveillance application. In our scenario it was determined that the most optimal position of the camera would be the ceiling, making the camera point down to the monitored area. This would provide us with an increased field of view, and would make detection of multiple people walking side by side easier. However, due to resource limitations and difficulties we would face trying to position the camera in such way, decision was made to simply place the camera on a higher floor and point it down to the monitored area in approximately 50-60° angle.

Object Extraction. As we have already touched a bit in Related Work section, background subtraction plays an important role in detecting people and tracking their movement. There are multiple ways in dealing with background subtraction, but we have mainly covered two - simple background subtraction - where we use a static background image to detect difference in frames - and running average approach - where we gradually create an approximate background image to extract moving objects from the frame. The former approach can work rather poorly and inaccurately in dynamic environments, whereas the latter one deals with the same challenges in a better way, thus have been chosen for our design. This is further discussed in greater detail in Section 3.3.1.

Object Detection. To build up a complete and reliable surveillance system, one must be able to identify whether the detected object is a person, an animal or simply tree leaves moving because of wind. To deal with this, one can possibly apply face detection methods or object's shape recognition techniques. Since these are rather complicating areas in image processing, and our test scenarios mainly dealt with environments where the background is more stable, we did not consider object identification and simply interpreted all changes in the monitored area as human movement. Our object detection approach is explained in more detail in Section 3.3.2.

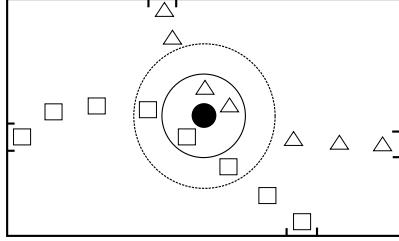


Figure 2.1: An illustration of the *k-nearest neighbors algorithm*.

Object Differentiation. Another important part of surveillance is being able to differentiation between multiple people crossing the monitored area. To deal with this challenge, we looked into two approaches, namely histogram approach - where we calculate individual histogram for every person and try to match it next time we detect people in the frame - and last coordinate approach - where we, instead of calculating person's histogram, calculate person's last coordinates in the monitored area. The main problem these both approaches share is that they are not very precise, however, the latter approach was easier to implement, thus was chosen for our design. Further comparison and discussion on object differentiation is provided in Section 3.3.3.

2.2 Prediction

In our scenario where occupants enter and exit the room constantly, typically within seconds or minutes of each other, it is somewhat difficult to find the immediate use of predicting an occupant's future actions. However, our partner team from Kenya expressed a need for such a system because they have had issues with thefts. Incorporating such a feature to their surveillance systems could warn a guard whenever suspicious activity might occur. A suspicious activity could involve a person moving too close to a given exit - possibly leading to a restricted area. In order to calculate a somewhat reliable prediction, we need to capture and store data about the behavior of previous occupants. When the data set grows, the prediction gradually becomes smarter and generally more precise. In order to satisfy the prediction requirement, various existing solutions have been considered, namely the *k-nearest neighbors algorithm* and an existing prediction model, the hidden Markov model. The two solutions will be covered briefly. The former assigns a given object to a group depending on the k nearest objects. To apply this to our scenario we imagined previously observed occupant positions being separated into groups dependent upon the chosen exit. Given any position we would analyze the k neighboring previous positions of distinct occupants and predict the exit to be that of the majority. Figure 2.1 depicts an example where the center dot is the current position and the squares and triangles are previous occupant positions who chose separate exits. The circle with the solid line resembles a situation where $k=3$ which indicates that the occupant at the center is most likely to take the same exit as the

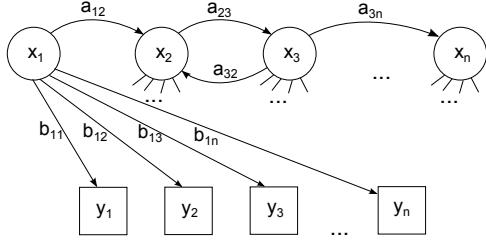


Figure 2.2: An illustration of the hidden Markov model.

occupants previously positioned at the triangle positions. The circle with the dashed line resembles a situation where $k=5$. This yields a different output where the occupant is most likely to choose the same exit as the occupants previously residing at the positions of the squares. The approach has proven to be inefficient on larger data sets [5], which can be averted by performing clever data reduction. This, however, deemed the approach out of the scope of this project.

Figure 2.2 depicts the different parameters a hidden Markov model works with. Each x is a state and each y is a possible observation, where a is a state transition probability and b is an output probability. When translating this to our scenario we need to define what exactly a state and an observation is and how to calculate the different probabilities. Since the ultimate goal is to predict what exit an occupant is going to take, a possible observation could be that an occupant exits a room at a given location. A state could depict a room area where a transition to another room area would have a state transition probability a . Given an occupant resides in any room area, b is the probability that the person will exit at a given location. Typically, the process is as follows; random probability values are assigned to a and b (maybe a ref to jahmm or something else). Gradually, these values are updated to reflect actual observed actions of occupants. Over time, the probabilities are modified and represent how the average occupant is moving given a current location, resulting in more accurate predictions - this process is commonly referred to as training. Even with the inclusion of a training procedure, the predictions will be incredibly unreliable in the beginning until a large enough data set has been collected. (DISCUSS, this might be wrong) Furthermore, imagine a scenario where an even number of occupants move from one exit to another in both directions, evenly split. At any location along the path the probabilities to each exit remain equal, since equally many previous occupants have taken each exit given the current location. In an attempt to partially avoid these issues, or at least produce some more reliable results with lesser data sets, we wanted to establish rule sets and integrate those into the calculations. Thus, we chose to design our own prediction model, heavily inspired by the hidden Markov model, since the original interpretation of states and probabilities as depicted in figure 2.2 is maintained. The custom model is explained in detail in Section 3.5.2.

Figure 2.2 depicts the different parameters a hidden Markov model works with. Each x is a state and each y is a possible observation, where a is a state transition probability and b is an output probability. When translating this to our scenario we need to define what

exactly a state and an observation is and how to calculate the different probabilities. Since the ultimate goal is to predict what exit an occupant is going to take, a possible observation could be that an occupant exits a room at a given location. A state could depict a room area where a transition to another room area would have a state transition probability a . Given an occupant resides in any room area, b is the probability that the person will exit at a given location. Typically, the process is as follows; random probability values are assigned to a and b (maybe a ref to jahmm or something else). Gradually, these values are updated to reflect actual observed actions of occupants. Over time, the probabilities are modified and represent how the average occupant is moving given a current location, resulting in more accurate predictions - this process is commonly referred to as training. Even with the inclusion of a training procedure, the predictions will be incredibly unreliable in the beginning until a large enough dataset has been collected. (DISCUSS, this might be wrong) Furthermore, imagine a scenario where an even number of occupants move from one exit to another in both directions, evenly split. At any location along the path the probabilities to each exit remain equal, since equally many previous occupants have taken each exit given the current location. In an attempt to partially avoid these issues, or at least produce some more reliable results with lesser data sets, we wanted to establish rule sets and integrate those into the calculations. Thus, we chose to design our own prediction model, heavily inspired by the hidden Markov model, since the original interpretation of states and probabilities as depicted in figure 2.2 is maintained. The custom model is explained in detail in section 3.5.2.

2.3 Communication Model

As there is several devices which have to communicate between each other, we have considered two different solutions:

- Central Server, and
- P2p

To decide which solution suited this project the most, each possible solution has been considered from the following perspectives.

1. ...

2.3.1 Central Server

The main benefits of this solution is as follows.

- Central Storage

3 | Design

3.1 System Overview

Figure 3.1 displays the different components in the system and how they relate and communicate with each other.

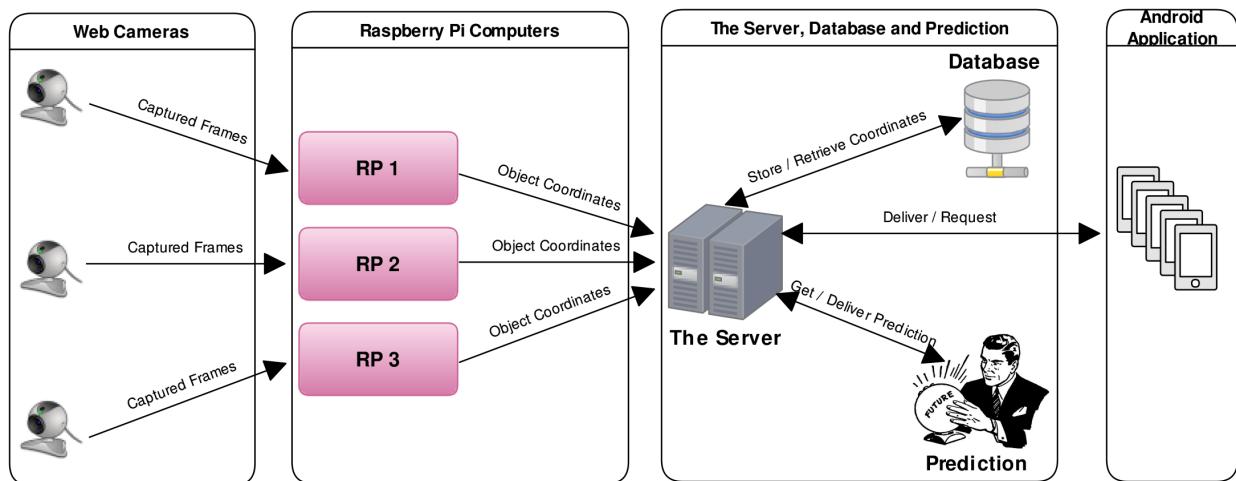


Figure 3.1: Occupancy Analyzer System Overview

As we can see, the system primarily consists of 4 components:

1. Web cameras,
2. Raspberry Pi computers,
3. The server,
4. and the Android application.

All of these components are discussed further on in the respective chapters.

3.2 Web Cameras

The purpose of the web cameras is simply to surveillance the area they have been placed in and forward the captured frames to Raspberry Pi computers for further processing, as shown in Figure 3.1. The cameras can be placed in a room, corridor, atrium or any other similar place in or outside of the building, where the services offered are required. Cameras can be placed either directly above the observed area (Figure 3.2) or in the corner of it, as illustrated in Figure 3.3. Naturally, a camera placed above the observed area would give better results, since this increases its field of view and makes it easier to correctly detect and distinguish between multiple people walking side by side. Furthermore, for the best results one must also take many different factors into account, such as the distance between the camera and monitored area, environmental conditions of the area the camera is placed in, lighting conditions.

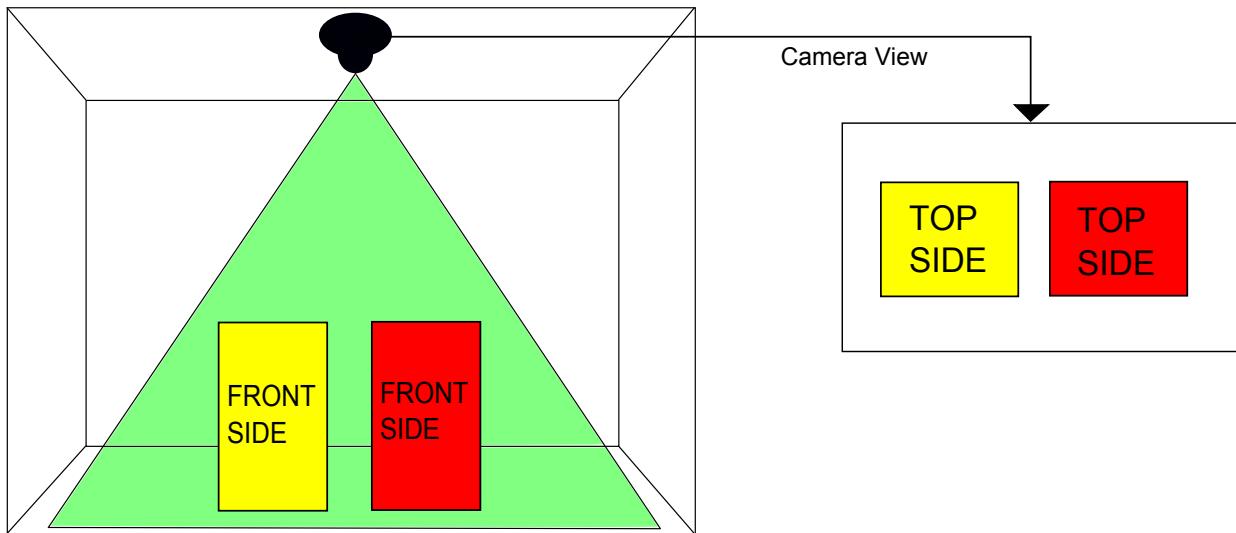


Figure 3.2: Camera top placement

3.3 Raspberry Pi Computers

The next component in the system architecture is Raspberry Pi computers. These computers have at least one web camera attached to them, and are responsible for processing the frames captured by these cameras. The main goal of processing these frames is to try to detect people in the monitored area and determine their position in that area. There are many different challenges in object detection, as well as various concepts and techniques that can be used to achieve this. We discuss these in the next sections.

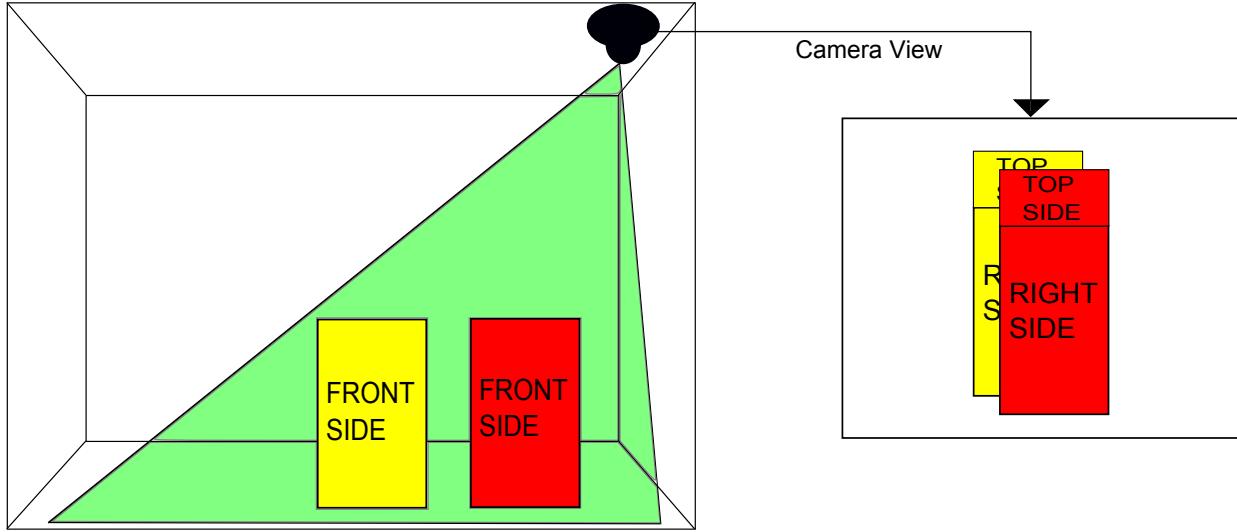


Figure 3.3: Camera corner placement

3.3.1 Object Extraction

To detect and extract objects, or in our case, people and their movement, we need to apply several motion detection techniques on the frames we are receiving from the web camera. First of all, to detect changes in some monitored area, we naturally need to have at minimum two images, which we must compare to see what changes occurred. We will try to look into two different approaches in doing this, a simple one, and one that is a bit more complicated and sophisticated, but much more adaptive and flexible.

3.3.1.1 Simple Background Subtraction Approach

A simple approach - called Background Subtraction - would be to have a static background image of observed area that was taken prior to the analysis and did not have any people in it. Then, one can simply detect changes and movement in the area by subtracting the static background image from every newly taken image of the monitored area [6]. The difference between the two images would then allow us to see if any changes happened, since after subtraction the resulting image would either be almost totally black (Figure 3.4) - meaning no one walked passed the observed area - or the image would have some resulting bright contours of detected differences (Figure 3.5).

After doing some research and experimenting with background subtraction technique, one will quickly discover that there are multiple weaknesses to it.

- First of all, if the initial background image is always static and never changes, this technique will fail in environments where lighting is dynamic. This is perfectly illus-



Figure 3.4: Background Subtraction with no background changes



Figure 3.5: Background Subtraction with background change

trated in Figure 3.6. We can see that the lighting is much darker in the second image, possibly because the light was turned off in the monitored room, thus after subtracting our static background image from this image, the resulting image is simply a lighter version of the two images, and not the intended black image. The reason why this happens is because the original brighter background image has a much higher intensity, thus the average value of its pixels are higher than the pixel values of the newly taken darker image. Naturally, this is a big problem, because now even if a person moves through the monitored area (Figure 3.7), he or she will not be as easily extractable as in Figure 3.5.

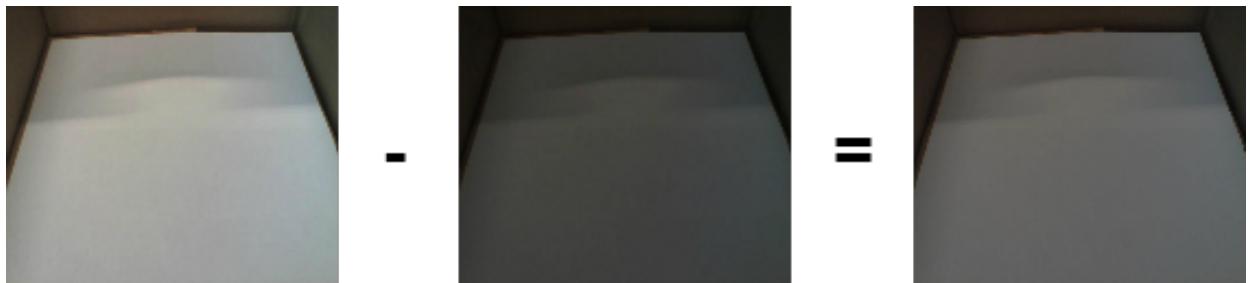


Figure 3.6: Background Subtraction with background lighting change

- Another problem with background subtraction approach surfaces when a camera is placed inside an area which has objects that constantly change their original position



Figure 3.7: Background Subtraction with background lighting change and lego figure appearing

(chairs, tables, appliances, etc.) by being moved, even small changes in object's location will spoil the resulting image after subtraction. As we can see in Figure 3.8, the object is displayed twice in the resulting image, even though we were not even interested in it, making it much harder to detect actual people moving in the area. From now on, the resulting image after subtraction will always be corrupt unless the object is placed back in its original position.

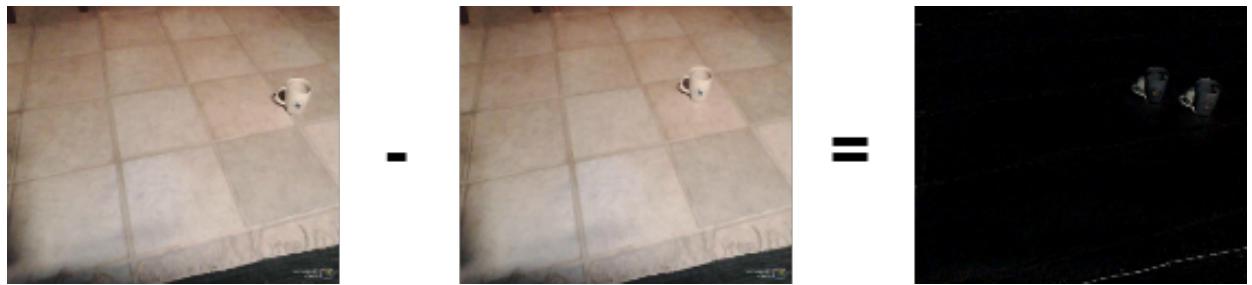


Figure 3.8: Background Subtraction with object changing it's location

In conclusion, we can see that simple background subtraction approach can work well in static environments, however it falls short in dynamic spaces. Naturally, these mentioned drawbacks of background subtraction approach need to be handled for object detection to work well, which creates additional challenges when implementing the system.

3.3.1.2 Running Average Approach

A much better approach for movement detection is using a running average method. In this technique we do not need to rely on a static background image of the monitored area taken prior to analysis. Instead, we try to find a new "approximate" background image by interpreting any changes in the background as noise and blurring them out. This is accomplished by taking a train sequence of multiple previously captured frames and performing arithmetic

averaging on that sequence [7]. This exact approach is illustrated in Figure 3.9. As we can see, hand motion moving up and down (Figures 3.9(a), 3.9(b) and 3.9(c)) gets blurred out when applying running average method, thus producing an approximate background image (Figure 3.9(d)) that can be used for subtraction of the test frames from it. This method of object extraction works very well and has a lot of flexibility. It can adapt to environmental changes in the monitored area, thus eliminating most of the weaknesses that the background subtraction approach has. For these reasons, the running average approach was chosen for our design.

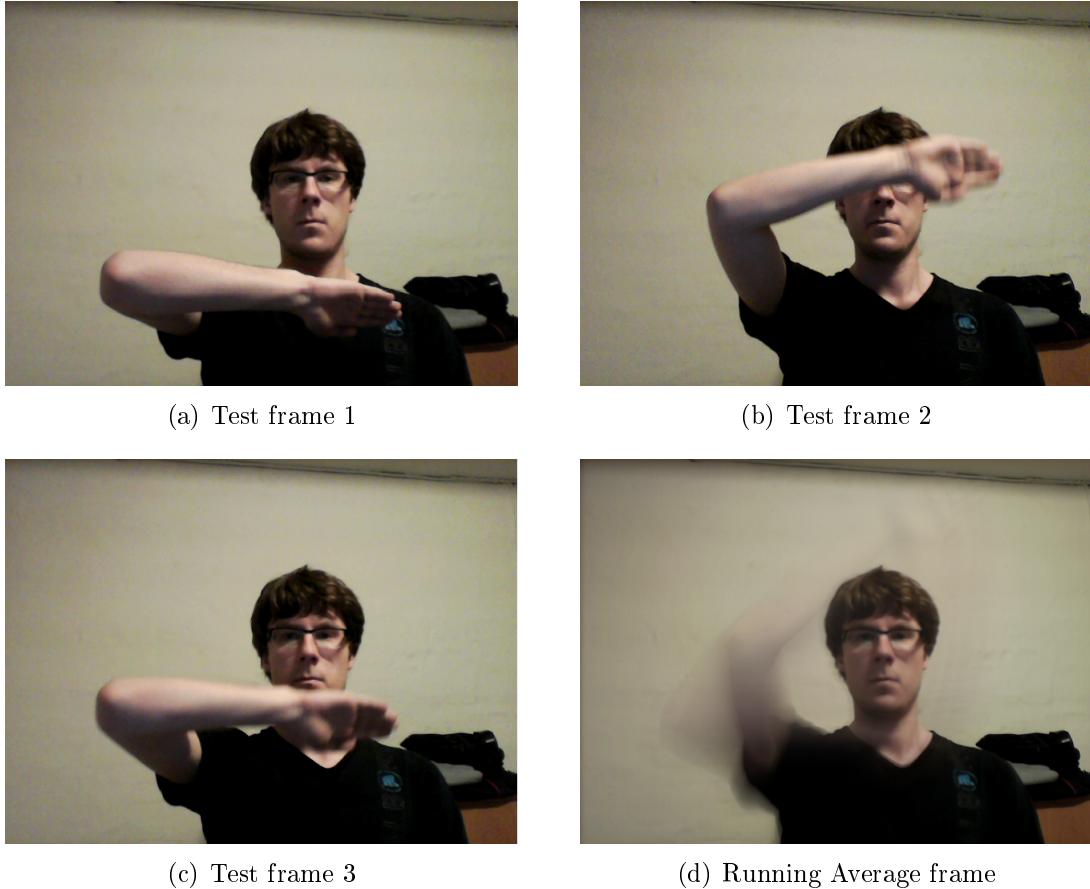


Figure 3.9: Example of Running Average technique

3.3.2 Object Detection

Now that we can extract objects using running average technique, we need to be able to actually find them in the resulting image we get after we perform subtraction. For this, we need to apply several key techniques in image processing.

- After we capture the initial frame of the monitored area, it will often contain noise and small details that we are not interested in. To deal with this, we must first apply blur or smoothing filter. In blurring technique we calculate weighted averages of areas of pixels in a source image by passing through it [8], which helps to reduce image noise and detail, as shown in Figure 3.10.

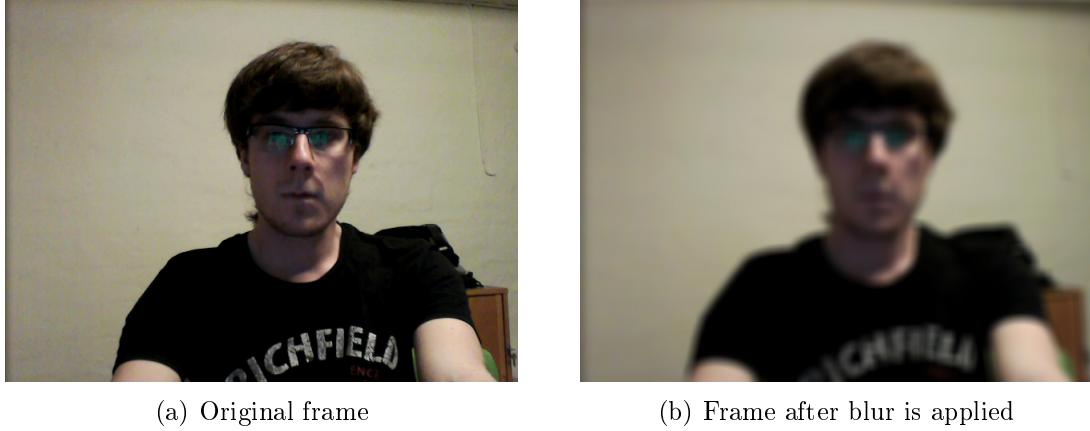


Figure 3.10: Blur example

- After we remove the initial noise, we can perform subtraction using running average approach (described in Section 3.3.1.2). After subtraction we will either get a nearly totally black image, meaning no motion occurred, or an image where some colors stand out, meaning some motion has occurred.
- In either case, for further processing of the taken frame, we can choose to convert it to a grayscale image. The reason for this is that the original RGB image we get has three channels, while grayscale image has only one [9], thus it is easier to work with. This procedure is illustrated in Figures 3.11(a), 3.11(b) and 3.11(c).
- For further processing of the image, we apply threshold technique, which converts the image to black and white and removes some more unwanted details and noise [10]. After threshold is applied we get the image shown in Figure 3.11(d).
- Moreover we want to expand the interesting parts of the image and contract smaller pieces, which can be considered noise and managed to slip through, even after we performed thresholding. To do so, we use two fundamental operations in morphological image processing, that is, dilation and erosion. Dilation allows us to expand the shapes contained in the image [11], whereas erosion simply shrinks shapes [12], so that bright regions surrounded by dark regions shrink in size, and dark regions surrounded by bright regions grow in size. When we apply dilation and erosion we get the image shown in Figure 3.11(e).

- Now, to project the detected area onto the original frame, we simply use bounding box technique, which gives us the coordinates of the rectangular border that fully covers the extracted white silhouette [13] that we got in 3.11(e). Then we use these coordinates to draw a simple rectangle around the extracted object, as well as mark its middle position by a red circle, as illustrated in Figure 3.11(f).

In conclusion, by applying the steps discussed in this section, we can fairly accurately detect people and their movement in the monitored area.

3.3.3 Object Differentiation

There will naturally be cases when multiple people will walk through the monitored area and will be captured by the cameras, therefore we must have a way to differentiate between them. This task becomes rather difficult if people are very close to each other, since they will simply be interpreted as one person. However, as long as people are far enough from each other, the task becomes significantly easier. There are multiple ways of differentiating between objects.

One of them is to simply look at the objects histogram, which gives a graphical representation of the intensity distribution of pixels [14]. Since people are usually dressed in different color clothes, we can simply calculate a histogram for every detected person and remember it. Now, every time we receive a new frame and detect a person in it, we go through our previously saved histograms and check whether any of them are the same as our newly detected persons histogram. If there is such histogram, we interpret the person we detected in our new frame as the same person we detected in the last, otherwise, we conclude that we have not detected this person before, thus save his histogram for future reference. The biggest weakness of this approach is that a person's clothes might have different colors from the front and back. Therefore, his histogram calculated while he is facing the camera might be rather different than the histogram of when his back was towards the camera. For this reason, if the person decides to turn around midway, he might be interpreted as a new person - never seen before by the camera - when in fact his frontal or back histogram was already saved.

Another approach of differentiating between multiple people, and in fact the approach we used in our design, is to simply use the whole frame as a coordinate system and remember the last coordinate of every single detected person. Now, similarly to histogram approach, whenever we detect a new person in the frame, we simply look throughout previously saved coordinates, and if we find that this new person's coordinates is relatively close to some previously saved person's coordinates, we simply interpret him as the same person we detected a second or few seconds ago, otherwise we see him as a new person. Naturally, we must regularly clear our previously saved coordinates, so that newly arrived person would not simply, by taking a similar path, be interpreted as a person who is no longer in the monitored area. This approach of differentiating between people is illustrated in Figure 3.12.

3.4 The Server

TODO

The server is the central part of the whole system. It is here the Raspberry PIs send all the gathered data, where data are stored and structured, and where Android devices can fetch information about rooms, people, and predictions on which exits a person is most likely to use.

The server is thus split into the following parts:

1. Data Input API,
2. Data Output API,
3. Storage.

3.4.1 Data Input API

All sorts of devices can be used to capture and process data. To simplify the process and to allow access for all types of devices, the server provides an API which allows the senders to focus on data capturing and processing.

3.4.2 Data Output API

As with the devices providing the data for the server, the devices which requests data from the server can be of different types. By providing an API these devices can focus on how to use and present the data, and let the server focus on how the data is fetched. This ensures that all data is available for everyone and that each request has a minimum of effect on other devices access to the data.

3.4.3 Storage

3.5 Prediction

3.5.1 Gathering statistical data

In addition to storing data about each occupant we also store data about how occupied each section of the recorded image is. The process is simple; we split the image of a room into

sections, referred to as cells, and each time an occupant enters a cell the stored activity for the given cell is increased. Each room has a different and independent set of cells. This data tells us which sections of the room are most occupied. We use the activity data to perform predictions about an occupant's future actions using a custom prediction model.

3.5.2 Custom prediction model based on HMM

We have chosen to build our own prediction model heavily inspired by the hidden Markov model. Each cell in the image is a state where the state transition probability is the likelihood of an occupant moving to an adjacent cell and the output probability is the likelihood of an occupant going to an exit given any current cell. Unlike a regular hidden Markov model, we do not store probability values individually for each state, but rather do the necessary calculations each time a prediction is requested by using the stored activity values of each cell. Additionally, our custom model allows us to take several custom factors into account during calculations. These factors work as a rule set for likely or unlikely occupant actions:

- An occupant entering a room from a given exit is less likely to exit the room at the given exit.
- An occupant is more likely to continue moving in his general direction and least likely to return to his previously visited cell.
- An occupant is more likely to move to the adjacent cell with the highest amount of previous activity.
- The likelihood of an occupant exiting at a given exit (unless it also serves as the occupant's entrance) is inversely proportional to the direct distance to the given exit, producing a magnet-like effect.

These factors have an influence on a final value of a cell or exit, which is used when calculating each probability. The sum of the probabilities of an occupant moving to each individual exit given a current cell is 100. The cell with the highest probability will be the predicted cell.

Figure 3.13 shows the custom model translated to a scenario. Each cell contains an identification number, an activity count and a value denoting whether or not the cell is an exit cell. The occupant entered the room at exit y_3 and the current state is x_{14} . a_{14-13} denotes the state transition probability of the occupant moving from state x_{14} to state x_{13} . b_{14-1} denotes the output probability of the occupant ultimately choosing exit y_1 . Taking our custom rule set into account, the occupant is least likely to exit at y_3 and most likely to continue to x_{13} . Moving to x_9 would increase the probability of the occupant exiting at y_1 (b_{9-1}) significantly.

3.6 Android Application

TODO



(a) Original frame



(b) After subtraction



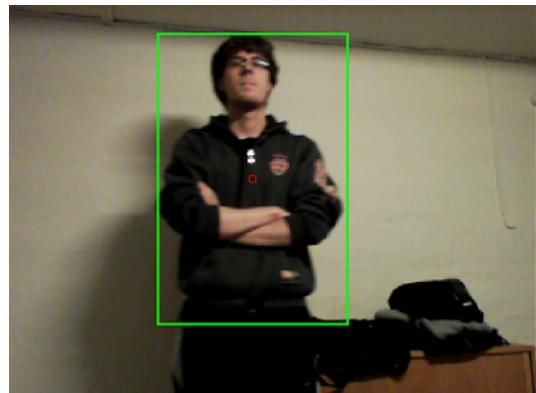
(c) After grayscale filter is applied



(d) After threshold is applied



(e) After dilation and erosion

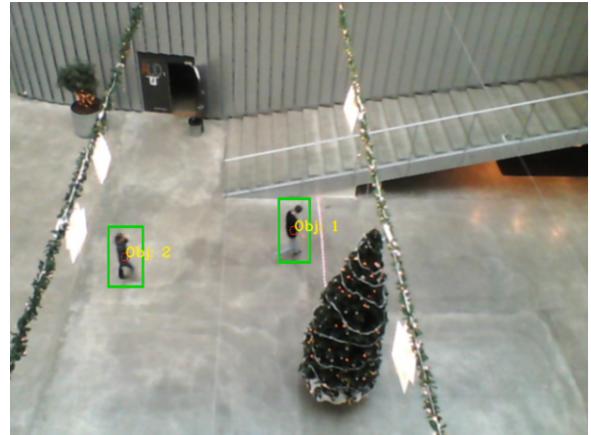


(f) After bounding box is drawn

Figure 3.11: Object detection example



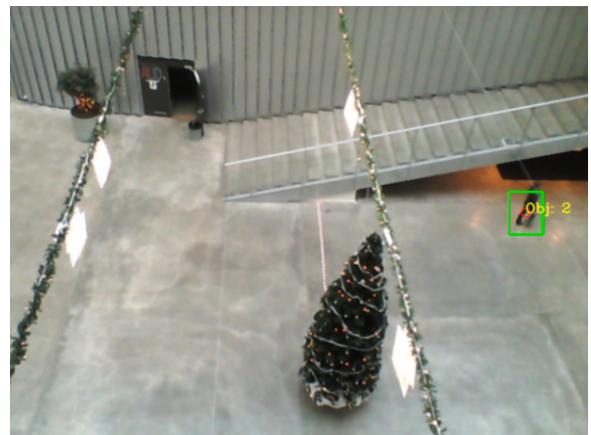
(a) One person enters the monitored area



(b) Second person enters the monitored area



(c) Two people continue their walk through the monitored area



(d) First person leaves the monitored area

Figure 3.12: Object differentiation example in the monitored area

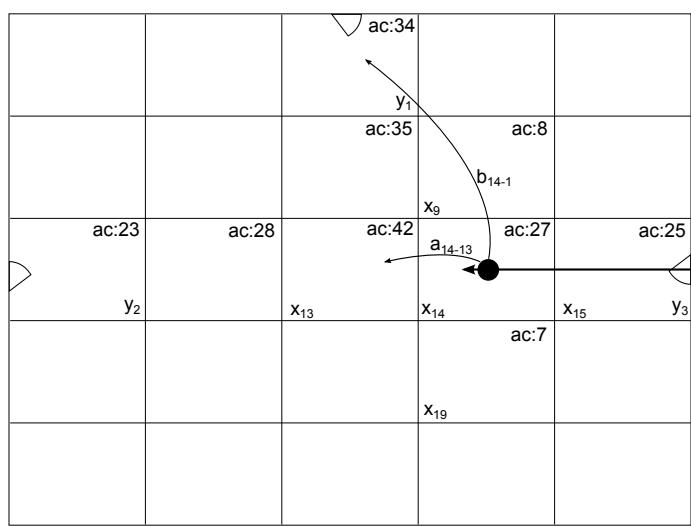


Figure 3.13: Applying the custom model to a scenario.

4 | Implementation

4.1 Image Processing

4.1.1 Library

As none of the group members had any previous experience with image processing, it was quite unrealistic to try to implement all of the image processing techniques needed for this project, and at the same time make them optimized enough, so it could work well and fast on Raspberry Pi computers with limited resources. For this reason the group decided to use an existing image processing library to ease the implementation process. After some research and experiments, we decided go with one of the most popular and well documented libraries called OpenCV¹. OpenCV is an open source computer vision and machine learning software library that has more than 2500 optimized algorithms for image processing. Furthermore, it provides interfaces to multiple popular programming languages, including C++, C, Python and Java.

4.1.2 Programming Language

For determining which programming language to use for image processing on Raspberry Pi computers, skill and preference document was created, where both, ITU and Strathmore University, teams indicated their skill and preference for various programming languages. The two languages that stood out the most were Java and Python, thus to choose one of them, we decided benchmark² them in order to compare their performance. The results are presented in the next section.

¹OpenCV homepage: <http://opencv.org/>.

²Code used for benchmarks: Java - <http://itu.dk/people/tmis/javatest/>, Python - <http://itu.dk/people/tmis/pytest/>.

4.1.2.1 Performance Comparison Between Java and Python

Since the project dealt with a real-time vision application that had to process large amount of frames, we were interested in how fast Java and Python can perform different image processing algorithms. Benchmarks were performed on a laptop ³ and a Raspberry Pi computer ⁴ to give some perspective on how much faster a modern laptop is compared to a Raspberry Pi computer.

- At first we tried to benchmark how fast can Java and Python perform a simple matrix multiplication of two 300×300 size matrices. The results are illustrated in Figure 4.1. We can clearly see that Java was way faster than Python in this benchmark. It took Java less than half of a second to perform the multiplication of two 300×300 size matrices on a laptop, whereas it took more than 3 seconds to do the same in Python (Figure 4.1(a)). Multiplication on a Raspberry Pi computer (Figure 4.1(b)) was naturally much slower than on a laptop. Java was again much faster than its counterpart by dealing with the task in less than 23 seconds, whereas Python was very close to hitting 2 minute mark to accomplish the same task.

In conclusion, as it was expected Java convincingly won the first benchmark.

- The second benchmark involved testing how fast Java and Python can perform different image processing algorithms. For this benchmark we used the OpenCV library that we introduced earlier. OpenCV has Java and Python interfaces and all the computations are performed on the native level ⁵, hence we have some overhead that is equal to a cost of one or several API calls. Consequently, the purpose of this benchmark was to see, which language - Java or Python - has less overhead and can perform API calls faster.

As we can see in Figure 4.2, the difference between Java and Python in this benchmark was rather small on both, laptop 4.2(a) and Raspberry Pi computer 4.2(b). This was rather surprising at first, since Java had a big edge in the first benchmark. One can speculate that Python's interface to OpenCV library has been implemented in a more efficient and optimized way than Java's, thus the outcome. In general, the difference was only in terms of a few milliseconds, however, it was a good enough reason for us to choose Python's interface to OpenCV library for image processing part in this project.

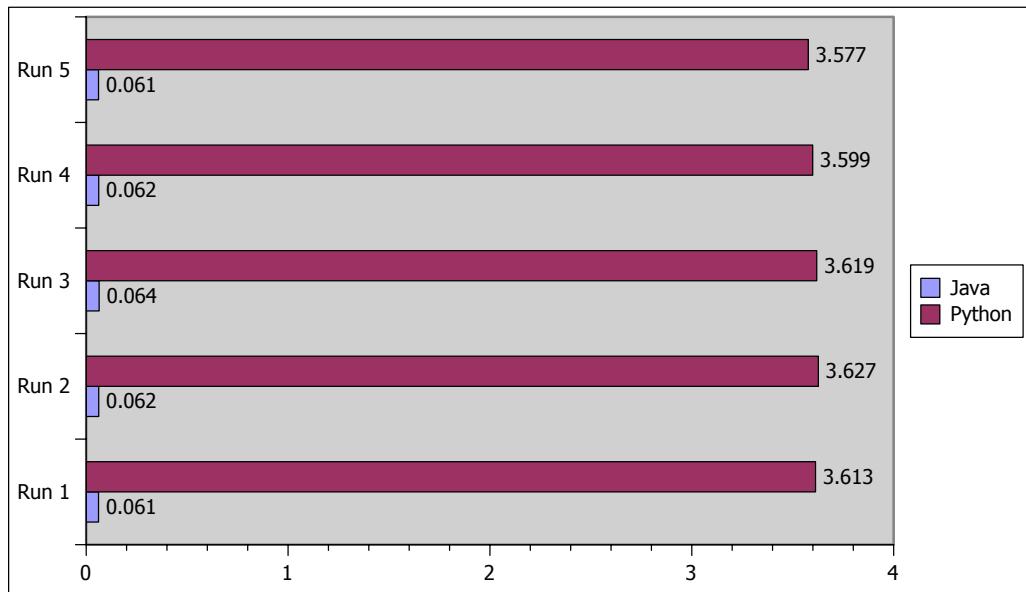
³Toshiba Satellite L855 Laptop (Intel Core i7-3630QM 2.40 GHZ, 4GB DDR3 1600MHz, Radeon HD7670M 2GB, Windows 7 OS).

⁴Raspberry Pi Computer (ARM1176JZF-S 700MHz, 512 MB memory, Broadcom VideoCore IV graphics, Linux Raspbian OS).

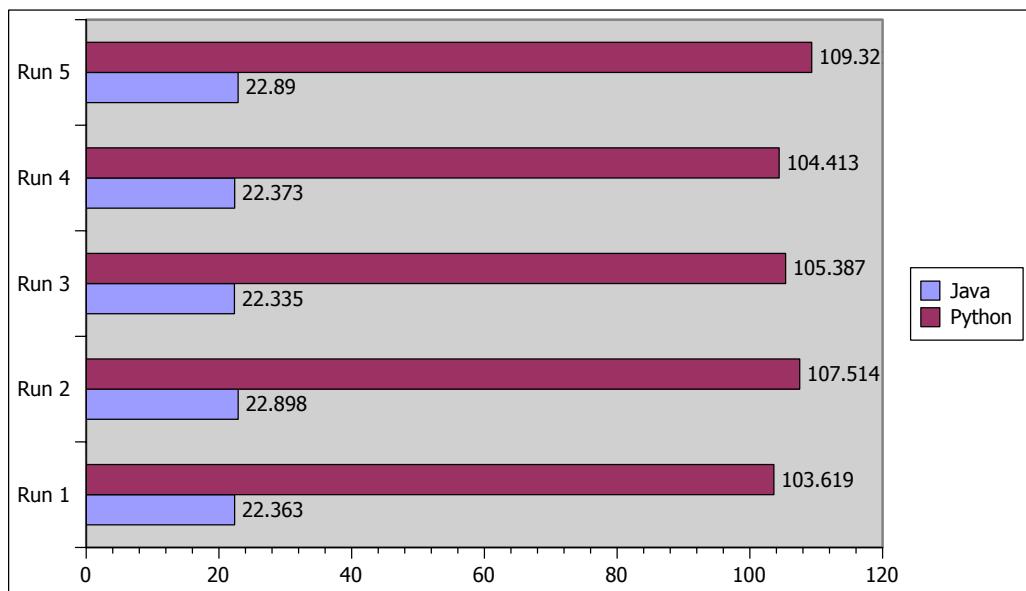
⁵In C++.

4.2 The Server

Server...

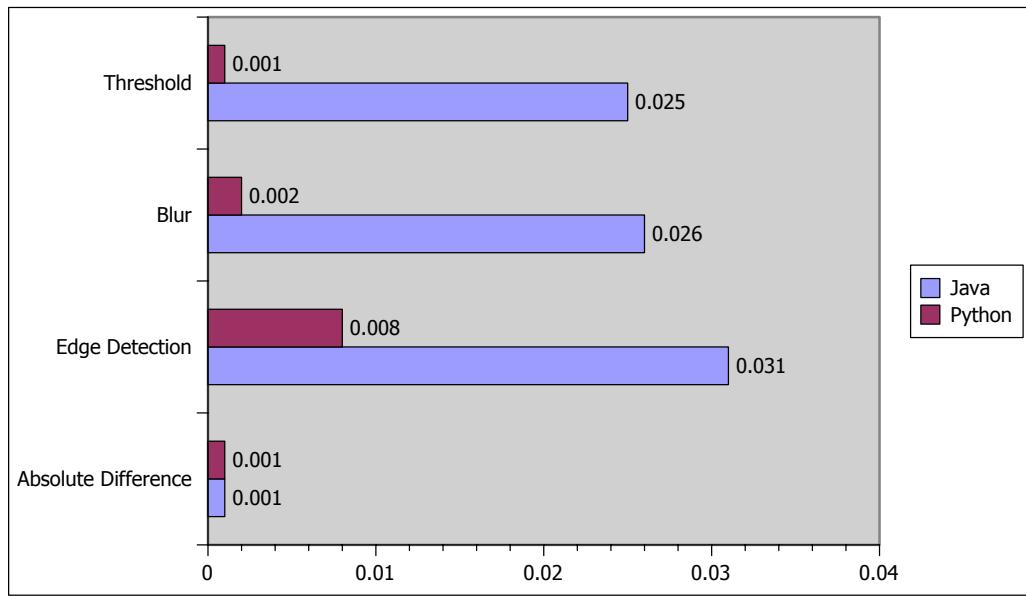


(a) Time (in seconds) needed to multiply two 300×300 size matrices in Java and Python on a laptop

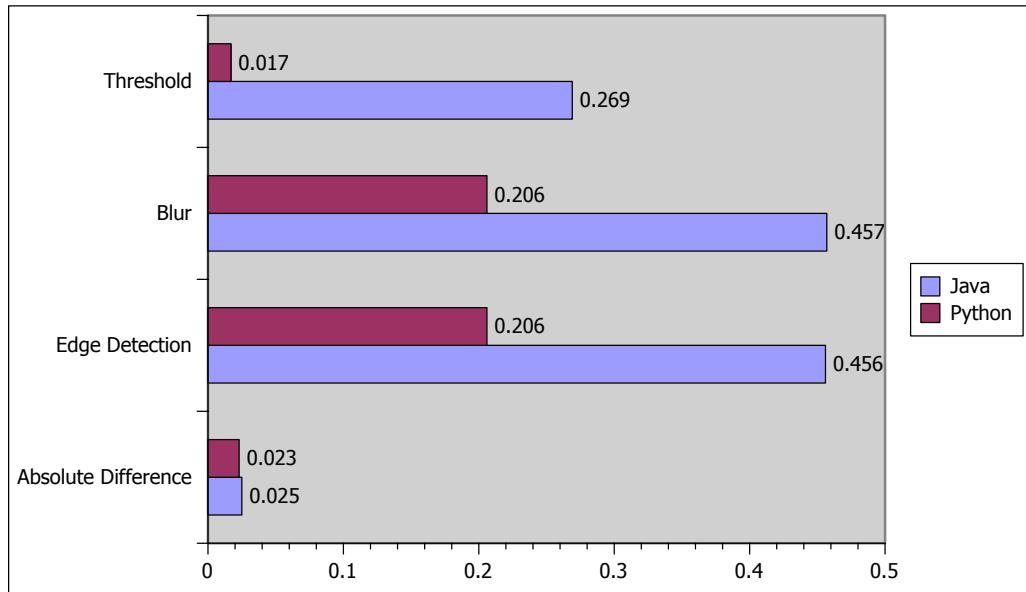


(b) Time (in seconds) needed to multiply two 300×300 size matrices in Java and Python on a Raspberry Pi computer

Figure 4.1: Benchmark of matrix multiplication of two 300×300 size matrices



(a) Time (in seconds) needed to process images in Java and Python on a laptop



(b) Time (in seconds) needed to process images in Java and Python on a Raspberry Pi computer

Figure 4.2: Benchmark of image processing

5 | Evaluation

5.1 Verification

5.2 Benchmarks

6 | Collaboration

This chapter deals with the collaboration and process in the project. First we introduce some theory on project management and global collaboration. In the next step we report our project work: We introduce the project team, the initial situation and present the chosen project management approach, methods and tools. In addition we evaluate the collaboration tools we used in the project. Finally, we lead to the issues we had to face during the project. These include progress issues, issues with the used project management techniques and collaboration issues. We analyze the collaboration issues based on the previous introduced theory. Furthermore we reflect the process and collaboration and explain our learning outcomes.

6.1 Theory

6.1.1 Project management

There are two major approaches how you can perform project management: The traditional approach and the agile approach.

The traditional approach is ...

_____ TODO

An agile project management approach is ...

_____ TODO

A popular agile project management framework is Scrum.

_____ TODO

6.1.2 Global collaboration

_____ TODO

- Trust
- Individualism

- Power distance

6.2 Background information of the Project Team

The project team consists of two groups of students. One group is from the Strathmore University located in Nairobi (Kenya). The other group is from the IT University (ITU) in Copenhagen (Denmark). The project team agreed on to name the two groups "Team Kenya" for the student group from Strathmore University and "Team ITU" for the student group from ITU. This helped to address each group in meeting reports, emails and conversations.

In the following paragraphs the background information for this project from both groups are introduced. The information explains the motivation and contribution of each group in the project and is used as argumentation why certain decisions were made and issues aroused.

6.2.1 Team ITU

Team ITU started with four members, which are all in the Masters programme "Software Development and Technology (Software Engineering)" of the ITU.¹ The members are from three different countries: Lithuania, Germany and Denmark.² The communication language within Team ITU is English, which is not the mother tongue of any of the members. The members of Team ITU met each other the first time on the 27.8.2013.

The students of Team ITU have to complete the project under the course "Global Software Development Project", which is mandatory for their masters programme. The requirements and deadlines for the project are given in the course base from ITU (reference) and by the advisor of the project. The course is rated with 15 ECTS points, which corresponds to approximately 20 hours per week per student. The students of Team ITU have to hand in this report as a mandatory requirement.

Link
to the
course
base

As the course for the Team ITU started in late August and the project team and topic was already known, Team ITU already started with the project work before Team Kenya. Team ITU and their advisor did not know when to get the contact details from the student group in Kenya, so the advisor recommended already initializing the project and doing some research and thoughts on the project.

¹In the beginning there were five members, but one left the project after two weeks, because he changed to another project and project team. This did not have an impact on the work as it was in the early stage of the project.

²Although there are minor differences between the nationalities, which could have an influence on the team work, we will not go into this, because it is out of scope for this report.

At the beginning Team ITU had received a different project topic by the advisor. The topic of the first project was "Vector Shooter".³ Team ITU put some thoughts into the project and spent time on defining the program, which they wanted to develop.

After three weeks the advisor had to change the topic of the project due to the collaboration with the Strathmore University. Only one student from the Strathmore University was interested on the project "Vector Shooter". Team ITU could have spent time and effort in convincing the other students from Strathmore University to do the project "Vector Shooter". Despite frustration, in consideration of the given deadline of the course and on recommendation of the advisor, Team ITU decided to not take this option.

Team ITU received the project topic and the contact details of the members of Team Kenya on the 17.9.2013. So the collaboration between the student groups did not start before this date.

6.2.2 Team Kenya

In the beginning Team Kenya consists of three members, which are all in the Masters programme "Telecommunication and Innovation" of the Strathmore University. The members are all from Kenya. The official languages in Kenya are English and Swahili. The members of Team Kenya met each other the first time on the 1.10.2013 (Global Meeting Report 1.10.2013/todoAdd refernce). One member had to leave the group in the last third of the project due to workload of other projects and obligations (Global Meeting Report 26.11.2013/todoAdd refernce).

The department, which is responsible for the masters programme "Telecommunication and Innovation", wants to convince the school management from Strathmore University to invest into the technology and idea of the project to improve environmental conservation. Therefore this collaboration project was initiated. (Global Meeting Report 3.12.2013/todoAdd refernce) .

For the students from Team Kenya the project is not included in a mandatory course and is completely voluntarily. The students from Team Kenya spend their free time on this project. There are no mandatory requirements or deadlines, which the Kenyan students have to achieve, except that they have to create documentation for their advisor to prove the progress of the project (Global Meeting Report 19.11.2013 /todoAdd refernce). The motivation for the Kenyan students to participate in this project is on the one hand to support the environmental thought and on the other hand to use this project work as basis

³In this project a program had to be developed, in which the calculation of a vector - based on an image - has to be made. The image should contain a person, who uses his hands to demonstrate a vector by taking a position like he uses a bow and arrow. The requirement was to use webcams, to capture the image, and to use RaspberryPis to detect the hands of the person and to calculate the vector. Furthermore, an Android application should be included in the program.

Maybe
in the
Introduction
Context
Part?

for their master thesis (Email 10.12.2013 /todoAdd refernce).

6.3 Project Management

6.3.1 Approach

As the course for the team ITU started in late August and the project team was already known, Team ITU already started with a project management technique, which was already established when Team Kenya got into the project.

Team ITU chose an agile project management approach, because it is more flexible, which was important due to the fact, that Team ITU did not knew the student group from Strathmore University neither their skills nor their requirements on this project. Moreover an agile project management approach is known for being more suitable for software development projects (*TODO: reference*).

Team ITU planned to merge Team Kenya into their project management approach as they did not come up with a different approach. Team ITU considered introducing Scrum to the project. However, due to the inexperience and non-knowledge about this framework in both teams, this idea has been dropped. For Scrum it is necessary that every project team member knows the concept. Thus, each team member would need to learn Scrum, which in turn would have meant losses for the implementation.

Find
refer-
ence

6.3.2 Project Organization

6.3.2.1 Timeline

Team Kenya agreed on to go with the deadlines from Team ITU as they do not have to meet any deadlines. (Global Meeting Report 24.11.2013 /todoAdd refernce)

6.3.2.2 Project Team and Roles

Each member of the project team created a member profile to introduce themselves, which are attached in the Appendix xx.

The members of Team ITU split up the tasks, such that different kind of roles aroused. The roles were not set in the beginning, but emerged during the project progress. They explain the major responsibilities. This does not

Add
refer-
ence

Collaboration master

Server Developer

Image Processing Developer

Prediction Model Developer

Team Kenya/Android Application Developers

6.3.3 Project Management Tools and Methods

6.3.3.1 Member profiles

Team ITU suggested to come up with some member profiles of each project team member to introduce themselves in the beginning of the project. One member of Team Kenya followed the suggestion in the beginning. The other two members provided their member profile three weeks after Team ITU requested them. It could be noted that this already led to disappointment on the side of Team ITU.

6.3.3.2 Skill-/Preference- and Motivation-Sheet

To see where the strengths, weaknesses and motivations within the project team are Team ITU came up with a Skill-/Preference-Sheet and Motivation-Sheet. Every project team member was encouraged to fill out these tables. The results can be seen in the Appendix xx (*reference*)/todoAdd refernce.

The outcome of these tables also influenced some decisions in the project. For example on which programming languages Team ITU made a benchmark on (*reference*)/todoAdd refernce, which database technology and which version control system for the source code is used. It also pointed out the possible roles of each project team member within the project (*reference*)/todoAdd refernce.

In the result it can be noticed that image processing, prediction models and RaspberryPis were quite unknown amongst the project team members. So each project team member was encouraged to do some research on these topics in the beginning.

6.3.3.3 "Out-of-Office"-Calendar

Because every project team member had different obligations besides the project, Team ITU came up with an "Out of Office"-Calendar, in which every project team member could publish for example exams, other hand-ins or other unavailability. The motivation of this calendar was to do a better planning on meetings and assignments during the project. The problem was that no one really looked into this calendar except in the beginning of the

project or when a certain issue aroused. The "Out-of-Office"-Calendar can be seen in the Appendix xx (*reference*)/todoAdd reference.

So it was noticed that Team Kenya as well as Team ITU had some exams in October. This influenced the project plan according to when to start the implementation phase. Later in the project it was forgotten or not correctly interpreted by Team ITU, that Team Kenya had a second exam phase in the first half of November. Team ITU got informed by this, when Team Kenya did not come up with agreed assignments.

6.3.3.4 Milestone plan

Team ITU chooses to create a very rough milestone plan to have an eye on the deadline given by ITU. A detailed milestone plan would contradict the idea of the agile project management approach. Team ITU did not know Team Kenya, their skills and requirements in the beginning and a detailed milestone plan probably would have been run over several times. In the milestone plan Team ITU roughly pointed out when they want to be done with each phase of the project.

The following milestones, where set in the beginning of the project.

Milestone	Description	Planned	Achieved
First Global Meeting	Introduction session with Team Kenya to get to know each other and sharing our requirements and expectations of the project	01.10.2013	01.10.2013
Final definition of Requirements and Roles	Get all requirements together and splitting up the project tasks	31.10.2013	03.12.2013
Finishing Prototype	A complete Prototype of the product	01.12.2013	15.12.2013
Report	A complete final report of the project	13.12.2013	16.12.2013

Figure 6.1: Milestone Plan

Figure 6.2 shows the milestone plan, which was planned in the beginning of the project.

The milestone plan was mostly a guideline for Team ITU, because they had a deadline in contrast to Team Kenya. Team ITU shared the planned dates for half of the milestones with Team Kenya as they agreed to go with the deadlines from Team ITU. (Global Meeting Report 24.11.2013) These shared milestones were according to when to finish the prototype and when to hand-in the mandatory final report.

Add
refer-
ence

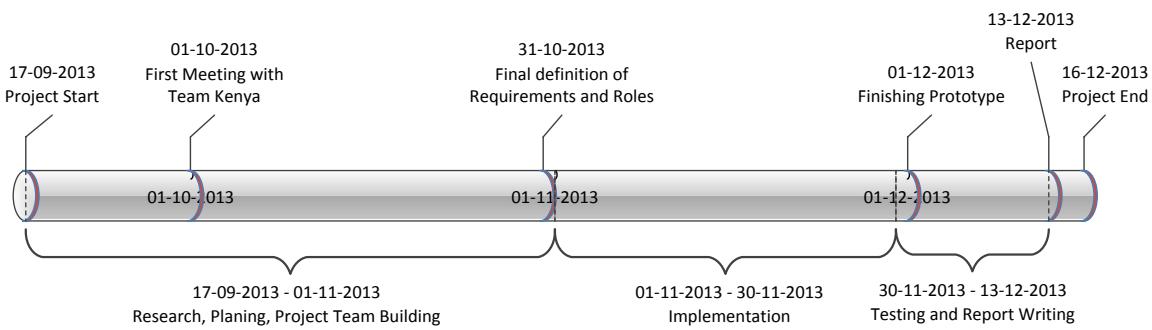


Figure 6.2: Milestone Timeline

6.3.3.5 Internal weekly meeting (Jour-Fixe)

Team ITU agreed on having an internal weekly meeting to update each other on the current progress, to discuss occurring topics, make decisions, work together and assign each other new assignments till the next meeting. Furthermore Team ITU also met with their advisor to check if they are on the right track. The concept of the weekly meeting equates to the concept of a Jour Fixe or a Daily Scrum.

There was a constant agenda for the meeting:

1. Status update of each team member
 - (a) What has each group done during the past week according to the project?
 - (b) Which achievements were made?
 - (c) Which problems occurred?
 - (d) Any other news
2. Discussions and Decisions (the topics in this section varied weekly according to the current progress)
3. Meeting with advisor
4. Assignments

This constant agenda gave some structure to the project. The team members know what to expect from the meeting and where they can place their concerns.

The members of Team ITU were mostly fully represented in the internal weekly meeting.

Find reference
Find reference

6.3.3.6 Global weekly meeting (Jour-Fixe)

Team Kenya and Team ITU agreed on having a weekly Skype-Meeting (*reference*) to update each other on the current progress, to discuss occurring topics, make decisions and assign each other new assignments till the next meeting, just like the weekly internal meeting. The meeting took place every Tuesday afternoon/evening. According to a Doodle-Survey (*reference*) all project team members could reserve this time for the project meeting.

Add reference

Each group could make a suggestion for the Agenda, but Team Kenya never used this opportunity. Therefore only Team ITU prepared an agenda, which was structured in the following way:

Add reference

1. Status update of each group
 - (a) What has each group done during the past week according to the project?
 - (b) Which achievements were made?
 - (c) Which problems occurred?
 - (d) Any other news
2. Discussions and Decisions (the topics in this section varied weekly according to the current progress)
3. Assignments

After the first few meetings Team Kenya and Team ITU agreed on to prepare the status update in advance, because it took quite a while for each group to present their status updates in the Skype chat. Unfortunately Team Kenya did not comply with this agreement except in the last meeting, where none of the Kenyans could attend due to other obligations or internet connection problems (Figure 6.3).

As it can been seen in Figure 6.3 almost every meeting took place. Team ITU was always in the majority. The attendance from Team Kenya was low mostly due to other obligations and internet connection problems.

6.3.3.7 Meeting reports

For every internal and global weekly meeting Team ITU wrote a meeting report. These meeting report documents all updates, news, discussions, decisions and assignments, which were discussed in the meetings. After publishing the meeting report, each project team member had the chance to correct wrong information or add missing information for the next couple of days. Afterwards the information and assignments in the meeting report were binding.

Global Meetings	Attendance Team ITU	Attendance Team Kenya	Prepared Status update Team ITU	Prepared Status update Team Kenya
01.10.2013	4	1	no	no
24.10.2013	4	2	no	no
31.10.2013	3	2	no	no
12.11.2013	4	2	no	no
19.11.2013	4	2	yes	no
26.11.2013	4	1	yes	no
03.12.2013	3	1	yes	no
10.12.2013	4	0	yes	yes

Figure 6.3: Milestone Plan

Project team members which could not attend the meeting had the chance to be updated with all necessary information. The report also confirms the content of the meeting. Furthermore the report is a source the team members can rely on. For example in case it comes to false allegations against a team member, the report could be used as a proof. The meeting report is also a reminder on the agreed tasks.

Team ITU expected that the agreed assignments were done or at least started to the next meeting. Since this did not happen on the side of Team Kenya, Team ITU started to put some deadlines on the most important assignments. But even those assignments were not done in time or done at all, although Team ITU got specific confirmation by email on the published meeting report or no disagreement on the meeting report from Team Kenya.

6.3.3.8 Time recording

6.3.3.9 Introduction and Kick-Off meeting

6.3.3.10 Surveys

6.3.4 Collaboration Tools

6.3.4.1 Communication

- Skype (excluding Google Hangout)
- Email

6.3.4.2 Sharing

- Google Drive (excluding Skydrive)
- Github

6.3.4.3 Time recording

- Toggl (excluding Excel)

6.4 Project Issues

6.4.1 Collaboration Issues

- Illusion of the project work and project team
- Failure to comply with the assignments
- Communication
- Lack of skills
- Other exams/hand-ins
- Differing requirements
- Attendence of meetings
- Equipment

6.5 Hypothetical Scenarios

- Assignments
- Communication
- Requirements
- Organisation by the universities (Requirements, clarification,)

7 | Discussion

8 | Conclusion

9 | References

- [1] Y. Benerezeth, P.-M. Jodoin, B. Emile, H. Laurent & C. Rosenberger. October 2012. *Comparative Study of Background Subtraction Algorithms*, pp. 5-6.
- [2] Aditi Job & Shirish Halbe. October 2012. *Multiple Objects Tracking using CAMshift Algorithm in Open CV*, pp. 41-46.
- [3] Gellert, A. & Vintan, L., 2006, *Person Movement Prediction Using Hidden Markov Models*, Computer Science Department, University of Sibiu, http://webspace.ulbsibiu.ro/arpad.gellert/html/SIC_HMM.pdf.
- [4] Ashbrook, D. & Starner, T., *Using GPS to Learn Significant Locations and Predict Movement Across Multiple Users*, Georgia Institute of Technology, <http://www.cc.gatech.edu/~thad/p/journal/using-gps-to-learn-significant-locations.pdf>.
- [5] Bhatia, N., 2010. *Survey of Nearest Neighbor Techniques*, p. 303. <http://arxiv.org/ftp/arxiv/papers/1007/1007.0085.pdf>
- [6] Massimo Piccardi. 2004. *Background subtraction techniques: a review*, p. 1.
- [7] Mohamad Hoseyn Sigari, Naser Mozayani & Hamid Reza Pourreza. February 2008. *Fuzzy Running Average and Fuzzy Background Subtraction: Concepts and Application*, p. 1.
- [8] Herman Cheng, Zhicong Huang & Mark Kumimoto. Spring 2006. *Final Project Report - Image Processing Techniques*, pp. 2-4.
- [9] Matt Siber. Fall 2005. *Resolution and File Size*, p. 1.
- [10] Brian Lovell *Thresholding Techniques*. Spring 2003.
- [11] R. Fisher, S. Perkins, A. Walker & E. Wolfart. 2003. *Dilation*.
- [12] R. Fisher, S. Perkins, A. Walker & E. Wolfart. 2003. *Erosion*.
- [13] Dr. Nicolas Pronost. *Introduction to Image Processing*, p. 9.
- [14] Ahmed Elgammal. Spring 2008. *Histograms of Digital Images*.

Todo list

TODO	32
TODO	32
TODO	32
TODO	32
Link to the course base	33
Maybe in the Introduction-Context Part?	34
Find reference	35
Add refernce	35
Add reference	37
Find reference	38
Find reference	38
Add reference	39
Add reference	39
Some note or other.	45
Another note.	45
And another one.	45
Figure: Add my picture here.	45

And another one.

Some
note
or
other.

Another
note.



Missing
figure

Add my picture here.